

Loading the Data and Importing Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

raw_2009_2010 = pd.read_excel('customer_transactions_sample.xlsx',
engine='openpyxl')
raw_2010_2011 = pd.read_excel('customer_transactions_sample.xlsx',
engine='openpyxl', sheet_name="Year 2010-2011")

raw_2009_2010.info()
raw_2010_2011.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 525461 entries, 0 to 525460
Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	Invoice	525461 non-null	object
1	StockCode	525461 non-null	object
2	Description	522533 non-null	object
3	Quantity	525461 non-null	int64
4	InvoiceDate	525461 non-null	datetime64[ns]
5	Price	525461 non-null	float64
6	Customer ID	417534 non-null	float64
7	Country	525461 non-null	object

dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 32.1+ MB

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541910 entries, 0 to 541909
Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	Invoice	541910 non-null	object
1	StockCode	541910 non-null	object
2	Description	540456 non-null	object
3	Quantity	541910 non-null	int64
4	InvoiceDate	541910 non-null	datetime64[ns]
5	Price	541910 non-null	float64
6	Customer ID	406830 non-null	float64
7	Country	541910 non-null	object

dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 33.1+ MB

```
raw_2010_2011.isna().sum() # total rows 541910
```

```

Invoice          0
StockCode        0
Description      1454
Quantity         0
InvoiceDate      0
Price           0
Customer ID     135080
Country         0
dtype: int64

raw_2009_2010.isna().sum() # total rows 525462

Invoice          0
StockCode        0
Description      2928
Quantity         0
InvoiceDate      0
Price           0
Customer ID     107927
Country         0
dtype: int64

```

Removing the rows with no description.

Why?

- Rows without descriptions lack this essential information, making it difficult to interpret or analyze the transactions accurately.
- Ambiguous data could lead to biased or misleading insights, affecting the validity of our conclusions.
- Removing rows with missing descriptions improves the overall quality of the dataset by eliminating incomplete or unreliable records. This enhances the effectiveness of subsequent analyses and modeling tasks that rely on accurate and complete data.

```

clean_2010_2011 = raw_2010_2011.dropna(subset=['Description'])
clean_2009_2010 = raw_2009_2010.dropna(subset=['Description'])

clean_2010_2011.isna().sum()
clean_2009_2010.isna().sum()

Invoice          0
StockCode        0
Description      0
Quantity         0
InvoiceDate      0
Price           0
Customer ID     104999

```

```
Country          0
dtype: int64
```

Merging the data from both 2009-2010 and 2010-2011 by adding a new Column in each dataframe for evaluation further.

```
clean_2009_2010['Year'] = '2009-2010'
clean_2010_2011['Year'] = '2010-2011'
clean_df = pd.concat([clean_2009_2010, clean_2010_2011],
ignore_index=True)
clean_df
```

Creating a categorical column for cancellations based on Invoice number

- The quantity for a cancellation is indicated by -ve numbers

```
clean_df['Cancellation'] = clean_df['Invoice'].apply(lambda x: 'Yes'
if str(x).startswith('C') else 'No')
clean_df.loc[clean_df.Cancellation=='Yes']

{"summary":{"\n  \"name\": \"clean_df\",\n  \"rows\": 19494,\n  \"fields\": [\n    {\n      \"column\": \"Invoice\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 8292,\n        \"samples\": [\n          \"C564940\",\n          \"C492712\",\n          \"C507284\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      {\n        \"column\": \"StockCode\",\n        \"properties\": {\n          \"dtype\": \"category\",\n          \"num_unique_values\": 2898,\n          \"samples\": [\n            20681,\n            22081,\n            22192\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n        {\n          \"column\": \"Description\",\n          \"properties\": {\n            \"dtype\": \"category\",\n            \"num_unique_values\": 3078,\n            \"samples\": [\n              \"FAWN BLUE HOT WATER BOTTLE\",\n              \"SET OF 60 I LOVE LONDON CAKE CASES\",\n              \"TEA BAG PLATE RED RETROSPOT\"\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n          },\n          {\n            \"column\": \"Quantity\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": 805,\n              \"min\": -80995,\n              \"max\": 1,\n              \"num_unique_values\": 207,\n              \"samples\": [\n                -94,\n                -7,\n                -500\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\"\n            },\n            {\n              \"column\": \"InvoiceDate\",\n              \"properties\": {\n                \"dtype\": \"date\",\n                \"min\": \"2009-12-01 10:33:00\",\n                \"max\": \"2011-12-09 11:58:00\",\n                \"num_unique_values\": 8141,\n
```



```

'rgb(255, 150, 238)', 'rgb(152, 115, 255)', 'rgb(151, 255, 145)',
'rgb(0, 225, 255)', 'rgb(255, 150, 140)', 'rgb(255, 196, 48)',
'rgb(255, 87, 101)', 'rgb(212, 255, 0)', 'rgb(166, 182, 255)',
'rgb(255, 136, 0)', 'rgb(255, 0, 162)', 'rgb(166, 0, 255)',
'rgb(216, 255, 209)', 'rgb(255, 194, 202)', 'rgb(250, 255, 110)'
]
# Group data by country and count the number of unique customers
customer_distribution = df.groupby('Country')['Customer
ID'].nunique().reset_index()
customer_distribution = customer_distribution.sort_values(by='Customer
ID', ascending=False).head(12)
fig = px.bar(customer_distribution, x='Country', y='Customer ID',
              title='Customer Distribution by Country',
              labels={'Country': 'Country', 'Customer ID': 'Number of
Customers'},
              color='Country', color_discrete_sequence=custom_palette)

fig.update_yaxes(range=[0, 1000])
fig.update_layout(
    title=dict(
        text='Customer Distribution by Country',
        x=0.5,
        font=dict(
            size=24
        )
    )
)
fig.show()

product_counts = df['Description'].value_counts().reset_index()
product_counts.columns = ['Description', 'Frequency']

product_counts = product_counts.sort_values(by='Frequency',
ascending=False)

top_products = product_counts.head(10)
fig = px.bar(top_products, x='Description', y='Frequency',
              title='Top 10 Most Frequently Purchased Products',
              labels={'Description': 'Product Description',
'Frequency': 'Frequency'},
              color='Description',
color_discrete_sequence=custom_palette)
fig.update_layout(xaxis_tickangle=-45)
fig.update_layout(
    title=dict(
        text='Top 10 Most Frequently Purchased Products',
        x=0.5,
        font=dict(

```

```

        size=24
    )
)
fig.show()

df['InvoiceDate'] = pd.to_datetime(df['InvoiceDate'])
df['YearMonth'] = df['InvoiceDate'].dt.to_period('M')
df['YearMonth'] = df['YearMonth'].astype(str)

# Group data by year and month and count the number of cancellations
cancellations_over_time = df[df['Cancellation'] ==
'Yes'].groupby('YearMonth').size().reset_index(name='Cancellation
Count')
fig = px.scatter(cancellations_over_time, x='YearMonth',
y='Cancellation Count',
                title='Frequency of Cancellations Over Time',
                labels={'YearMonth': 'Year-Month', 'Cancellation
Count': 'Cancellation Count'},
                size='Cancellation Count',
                size_max=12,
                )
fig.add_trace(px.line(cancellations_over_time, x='YearMonth',
y='Cancellation Count').data[0])

fig.update_layout(
    title=dict(
        text='Frequency of Cancellations Over Time',
        x=0.5,
        font=dict(
            size=24
        )
    )
)

fig.show()

revenue_over_time = df.groupby('YearMonth')
['Price'].sum().reset_index()
fig = px.scatter(revenue_over_time, x='YearMonth', y='Price',
                title='Revenue Trend Over Time',
                labels={'YearMonth': 'Year-Month', 'Price':
'Revenue'},
                size='Price',
                size_max=12,
                color_discrete_sequence=['springgreen']
                )
fig.add_trace(px.line(revenue_over_time, x='YearMonth',
y='Price', color_discrete_sequence=['limegreen'] ).data[0])
fig.update_layout(

```

```

        title=dict(
            text='Revenue Trend Over Time',
            x=0.5,
            font=dict(
                size=24
            )
        )
    )
)

fig.show()

# Convert 'InvoiceDate' to datetime format
df['InvoiceDate'] = pd.to_datetime(df['InvoiceDate'])
df['YearMonth'] = df['InvoiceDate'].dt.to_period('M')
df['YearMonth'] = df['YearMonth'].astype(str)

# Group data by customer and year-month, and calculate the number of
# active customers for each period
active_customers = df.groupby(['YearMonth', 'Customer
ID']).size().reset_index(name='Active')

active_customers['PreviousActive'] =
active_customers.groupby('Customer ID')['Active'].shift(1)

active_customers['Retained'] = active_customers['Active'] /
active_customers['PreviousActive']

active_customers['Retained'].fillna(0, inplace=True)

retention_rate = active_customers.groupby('YearMonth')
['Retained'].mean().reset_index()
fig = px.scatter(retention_rate, x='YearMonth', y='Retained',
                 title='Customer Retention Rate Over Time',
                 labels={'YearMonth': 'Year-Month', 'Retained':
'Retention Rate'},
                 size='Retained',
                 size_max=12,
                 )
fig.add_trace(px.line(retention_rate, x='YearMonth',
y='Retained').data[0])

fig.update_layout(
    title=dict(
        text='Customer Retention Rate Over Time',
        x=0.5,
        font=dict(
            size=24
        )
    )
)

```

```
    )  
)  
fig.show()
```